

# Milestone 1: Submission Checklist

MLOps Course - Module 2

## Pre-Submission Verification

Use this checklist to verify your submission before the deadline. Complete each section to ensure you meet all requirements.

---

## Deliverables Checklist

### FastAPI Service (Local)

- main.py exists with FastAPI app exposing /predict endpoint
- Pydantic request model defined for input validation
- Pydantic response model defined for output schema
- Model artifact (model.pkl or similar) included in repository
- Model loads deterministically at startup (not per-request)
- requirements.txt OR pyproject.toml with exact version pinning
- README describes lifecycle position (input → model → API → consumer)

### Cloud Run Deployment

- Cloud Run service URL is publicly accessible with HTTPS
- GCP Artifact Registry image reference documented
- Evidence of successful inference (screenshot or curl output)
- Cold start behavior analysis included in documentation

### Serverless Function (GCP Cloud Functions)

- Cloud Function code implementing prediction logic
- Deployment configuration documented
- Deployment logs captured (screenshot or text)
- Function invocation tested and working

### Comparative Report

- FastAPI container vs Cloud Function comparison included
- Lifecycle differences (stateful vs stateless) explained
- Artifact loading strategies compared
- Latency characteristics documented (cold starts, warm instances)
- Reproducibility considerations discussed

---

## Documentation

- README includes setup and deployment instructions
  - API usage examples with sample requests/responses
  - Lifecycle stage explanations present
  - Model-API interaction clearly described
  - Deployment URLs included and accessible
- 

## Rubric Evidence Map

Use this table to verify you have evidence for each graded criterion:

Criterion	Points	Evidence Location	Verified
Correct API implementation with Pydantic schemas	2	main.py - endpoint and model definitions	[ ]
Deterministic artifact loading	2	main.py - startup loading logic	[ ]
Reproducible environment	2	requirements.txt or pyproject.toml	[ ]
Successful HTTPS deployment on Cloud Run	2	Cloud Run URL + screenshot/curl output	[ ]
Proper registry workflow	2	Artifact Registry image reference	[ ]
Working Cloud Function inference	2	Function URL + invocation evidence	[ ]
Clear deployment stage explanation	2	README - Lifecycle section	[ ]
Artifact management documentation	2	README - Model-API interaction section	[ ]
Latency/cold start comparison	1	Comparative report section	[ ]
Statelessness/reproducibility comparison		Comparative report section	[ ]

---

Criterion	Points	Evidence Location	Verified
Clear instructions and organized code	1	Overall README and repository structure	[ ]

---

## Common Pitfalls

\*\* Warning:\*\* Hardcoding credentials or API keys

Never commit secrets to your repository. Use environment variables and GCP IAM:

```
# Bad
API_KEY = "sk-abc123..."

# Good
API_KEY = os.environ.get("API_KEY")
```

\*\* Warning:\*\* Model artifact not versioned or reproducibly loadable

Your model must load identically every time. Verify with:

```
# Ensure deterministic loading
model = joblib.load("model.pkl")
assert model.predict([[1, 2, 3]]) == model.predict([[1, 2, 3]])
```

\*\* Warning:\*\* Only testing warm instance latency

Cold starts behave differently. Test both:

```
# Wait 15+ minutes, then test cold start
curl -w "\nTime: ${time_total}s\n" $CLOUD_RUN_URL/predict
# Immediately test warm instance
curl -w "\nTime: ${time_total}s\n" $CLOUD_RUN_URL/predict
```

\*\* Warning:\*\* Loading model on every request instead of at startup

**Incorrect (slow):**

```
@app.post("/predict")
def predict(data: InputModel):
    model = joblib.load("model.pkl") # Loaded every request!
    return model.predict(...)
```

**Correct (fast):**

```
model = joblib.load("model.pkl") # Loaded once at startup

@app.post("/predict")
```

---

```
def predict(data: InputModel):
    return model.predict(...)
```

\*\* Warning:\*\* Cloud Run service not publicly accessible

Ensure you've configured public access:

```
gcloud run services add-iam-policy-binding SERVICE_NAME \
    --member="allUsers" \
    --role="roles/run.invoker"
```

\*\* Warning:\*\* Missing Pydantic validation schemas

Both request and response must use Pydantic models:

```
from pydantic import BaseModel

class PredictRequest(BaseModel):
    features: list[float]

class PredictResponse(BaseModel):
    prediction: float
    model_version: str
```

---

## Automated Sanity Checks

Run these commands locally before submitting:

### File Existence Checks

```
# Check required files exist
echo "=== Checking required files ==="
test -f main.py && echo "✓ main.py exists" || echo "✗ main.py missing"
test -f README.md && echo "✓ README.md exists" || echo "✗ README.md missing"
(test -f requirements.txt || test -f pyproject.toml) && echo "✓ Dependency
    file exists" || echo "✗ No requirements.txt or pyproject.toml"
ls *.pkl 2>/dev/null && echo "✓ Model artifact found" || echo "✗ No .pkl
    model artifact found"
```

### Dependency Pinning Validation

```
# Check for unpinned dependencies
echo "=== Checking for unpinned dependencies ==="
if [ -f requirements.txt ]; then
    grep -E '^[a-zA-Z]' requirements.txt | grep -v '==' | grep -v '^#' &&
        echo " Found unpinned dependencies above" || echo "✓ All
            dependencies appear pinned"
fi
```

---

## FastAPI Local Test

```
# Test FastAPI locally
echo "=== Testing FastAPI locally ==="
pip install -r requirements.txt
uvicorn main:app --host 0.0.0.0 --port 8000 &
sleep 3
curl -X POST "http://localhost:8000/predict" \
-H "Content-Type: application/json" \
-d '{"features": [1.0, 2.0, 3.0]}'
pkill -f uvicorn
```

## Pydantic Schema Verification

```
# Check for Pydantic models in main.py
echo "=== Checking for Pydantic schemas ==="
grep -E "class.*BaseModel" main.py && echo "✓ Pydantic models found" ||
echo "✗ No Pydantic BaseModel classes found"
grep -E "@app\.(post|get).*response_model" main.py && echo "✓ Response
model specified" || echo " Consider adding response_model to endpoint"
```

## Cloud Run Deployment Verification

```
# Test Cloud Run endpoint (replace with your URL)
echo "=== Testing Cloud Run deployment ==="
CLOUD_RUN_URL="https://your-service-abc123.run.app"
curl -s -o /dev/null -w "%{http_code}" "$CLOUD_RUN_URL/predict" \
-X POST \
-H "Content-Type: application/json" \
-d '{"features": [1.0, 2.0, 3.0]}' | \
grep -q "200" && echo "✓ Cloud Run returns 200" || echo "✗ Cloud Run not
responding correctly"
```

## Cloud Function Verification

```
# Test Cloud Function endpoint (replace with your URL)
echo "=== Testing Cloud Function ==="
FUNCTION_URL="https://us-central1-your-project.cloudfunctions.net/predict"
curl -s -o /dev/null -w "%{http_code}" "$FUNCTION_URL" \
-X POST \
-H "Content-Type: application/json" \
-d '{"features": [1.0, 2.0, 3.0]}' | \
grep -q "200" && echo "✓ Cloud Function returns 200" || echo "✗ Cloud
Function not responding correctly"
```

---

## Latency Comparison Test

```
# Compare cold start vs warm latency
echo "== Latency Comparison =="
echo "Cloud Run cold start (after 15min idle):"
curl -w "Total time: %{time_total}s\n" -s -o /dev/null \
  -X POST "$CLOUD_RUN_URL/predict" \
  -H "Content-Type: application/json" \
  -d '{"features": [1.0, 2.0, 3.0]}'

echo "Cloud Run warm request (immediate follow-up):"
curl -w "Total time: %{time_total}s\n" -s -o /dev/null \
  -X POST "$CLOUD_RUN_URL/predict" \
  -H "Content-Type: application/json" \
  -d '{"features": [1.0, 2.0, 3.0]}
```

---

## Self-Assessment Questions

Answer these questions honestly before submitting:

### Reproducibility

- “Can someone clone my repo and run the FastAPI service in under 5 minutes?”**
  - Test: Clone to a different directory and follow your own setup instructions
- “Are ALL my dependencies pinned to exact versions?”**
  - Check: no >=, <=, ~=, or ^ in your dependency file
- “Does my model load identically every time?”**
  - Test: Load model twice and verify predictions match

### Deployment

- “Is my Cloud Run service publicly accessible right now?”**
  - Test: Open the URL in an incognito browser window
- “Does my Cloud Function respond correctly?”**
  - Test: Invoke via curl from a different machine/network
- “Have I captured evidence of both deployments working?”**
  - Screenshots or curl output saved and included

### API Design

- “Do my Pydantic schemas validate input correctly?”**
  - Test: Send malformed JSON and verify you get a 422 error
- “Does my API return structured responses?”**
  - All responses should match the Pydantic response model

### Lifecycle Understanding

- 
- “Can I explain where my deployment fits in the ML lifecycle?”**
    - Should connect: data → training → artifact → API → consumer
  - “Have I documented monitoring touchpoints?”**
    - Where would you add logging, metrics, alerts?

## Comparative Analysis

- “Have I tested and documented cold start behavior for both patterns?”**
  - Cloud Run and Cloud Functions behave differently
- “Can I explain the trade-offs between stateful and stateless deployment?”**
  - When would you choose each pattern?
- “Have I compared reproducibility across deployment patterns?”**
  - Container vs function: which is more reproducible and why?

## Final Verification

- “Have I committed and pushed all my changes?”**

```
git status # Should show "nothing to commit, working tree clean"
```

- “Are all my deployment URLs included in the README?”**
  - Cloud Run URL and Cloud Function URL both documented
- “Is my repository accessible to the instructor?”**
  - Public repo: Anyone can view
  - Private repo: Check Settings → Collaborators

---

## Quick Reference: Expected Repository Structure

```
your-repo/
├── main.py                      # FastAPI application
├── model.pkl                     # Trained model artifact
├── requirements.txt              # OR pyproject.toml
├── Dockerfile                    # For Cloud Run deployment
└── cloud_function/
    ├── main.py                  # Cloud Function code
    └── requirements.txt         # Function entry point
        # Function dependencies
├── README.md                     # Documentation with:
    ├── Setup instructions
    ├── API usage examples
    ├── Deployment URLs
    ├── Lifecycle explanation
    └── Comparative analysis
└── screenshots/                  # Optional: deployment evidence
    ├── cloud_run_response.png
    └── cloud_function_logs.png
└── tests/                        # Optional: local tests
    └── test_api.py
```

---

---

## **Submission Confirmation**

Once all checks pass:

1. Verify your Cloud Run URL responds (format: `your-service.run.app/predict`)
2. Verify your Cloud Function URL responds (format: `region-project.cloudfunctions.net/predict`)
3. Confirm comparative analysis is included in README
4. Copy your repository URL (format: `github.com/YOUR_USERNAME/YOUR_REPO`)
5. Submit the URL via the course submission system